

# RLC Hybrid Intelligence Architecture

## From Automated Scripts to Intelligent Agents: A Migration Plan

**Version:** 0.1 (Draft)

**Date:** February 4, 2026

**Author:** Architecture planning session with Claude

---

### Executive Summary

This document outlines a phased approach to evolving RLC's analytical infrastructure from deterministic automation to a hybrid system that combines:

- **Statistical rigor** (econometric models, ML forecasting)
- **AI reasoning** (LLMs for synthesis, judgment, and content)
- **Agentic coordination** (goal-oriented agents working together)

The architecture preserves existing investments, adds capabilities incrementally, and includes clear decision points for hardware purchases.

---

## Part 1: Current State Assessment

### Existing Infrastructure

Component	Status	Notes
MSI Aegis ZS2 (RTX 5080, 16GB VRAM)	<input checked="" type="checkbox"/> Operational	Primary compute, runs 24/7
Lenovo Yoga 9	<input checked="" type="checkbox"/> Operational	Secondary/mobile
PostgreSQL Database	<input checked="" type="checkbox"/> Operational	Bronze/Silver/Gold layers
46 Agent Scripts	<input checked="" type="checkbox"/> Operational	Mostly deterministic
Power BI Dashboards	<input checked="" type="checkbox"/> Operational	Visualization layer
Notion "RLC OS"	<input checked="" type="checkbox"/> Operational	Knowledge management

## Current Agent Inventory (Estimated Categories)

### Data Collection Agents (~15 agents)

- └─ USDA PSD API
- └─ USDA NASS
- └─ USDA AMS
- └─ NOAA Weather
- └─ Census Trade Data
- └─ FGIS Export Inspections
- └─ Various market feeds

### Data Processing Agents (~12 agents)

- └─ Bronze → Silver transformations
- └─ Data validation
- └─ Schema normalization
- └─ Marketing year alignment
- └─ Time series alignment

### Analytics Agents (~8 agents)

- └─ Basis calculations
- └─ Spread monitoring
- └─ Seasonal patterns
- └─ Trend detection

### Output Agents (~6 agents)

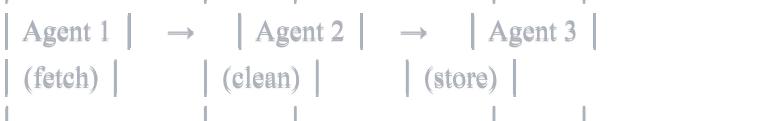
- └─ Report generation
- └─ Power BI refresh
- └─ Alert dispatch
- └─ Content formatting

### Utility Agents (~5 agents)

- └─ Error handling
- └─ Logging
- └─ Scheduling
- └─ Health checks

## Current Architecture Pattern

ORCHESTRATOR  
(Scheduled jobs, sequential execution)



Each agent: Single task, explicit rules, no reasoning

### Strengths of Current System:

- Reliable and predictable
- Fast execution
- Low resource usage
- You understand every line of code

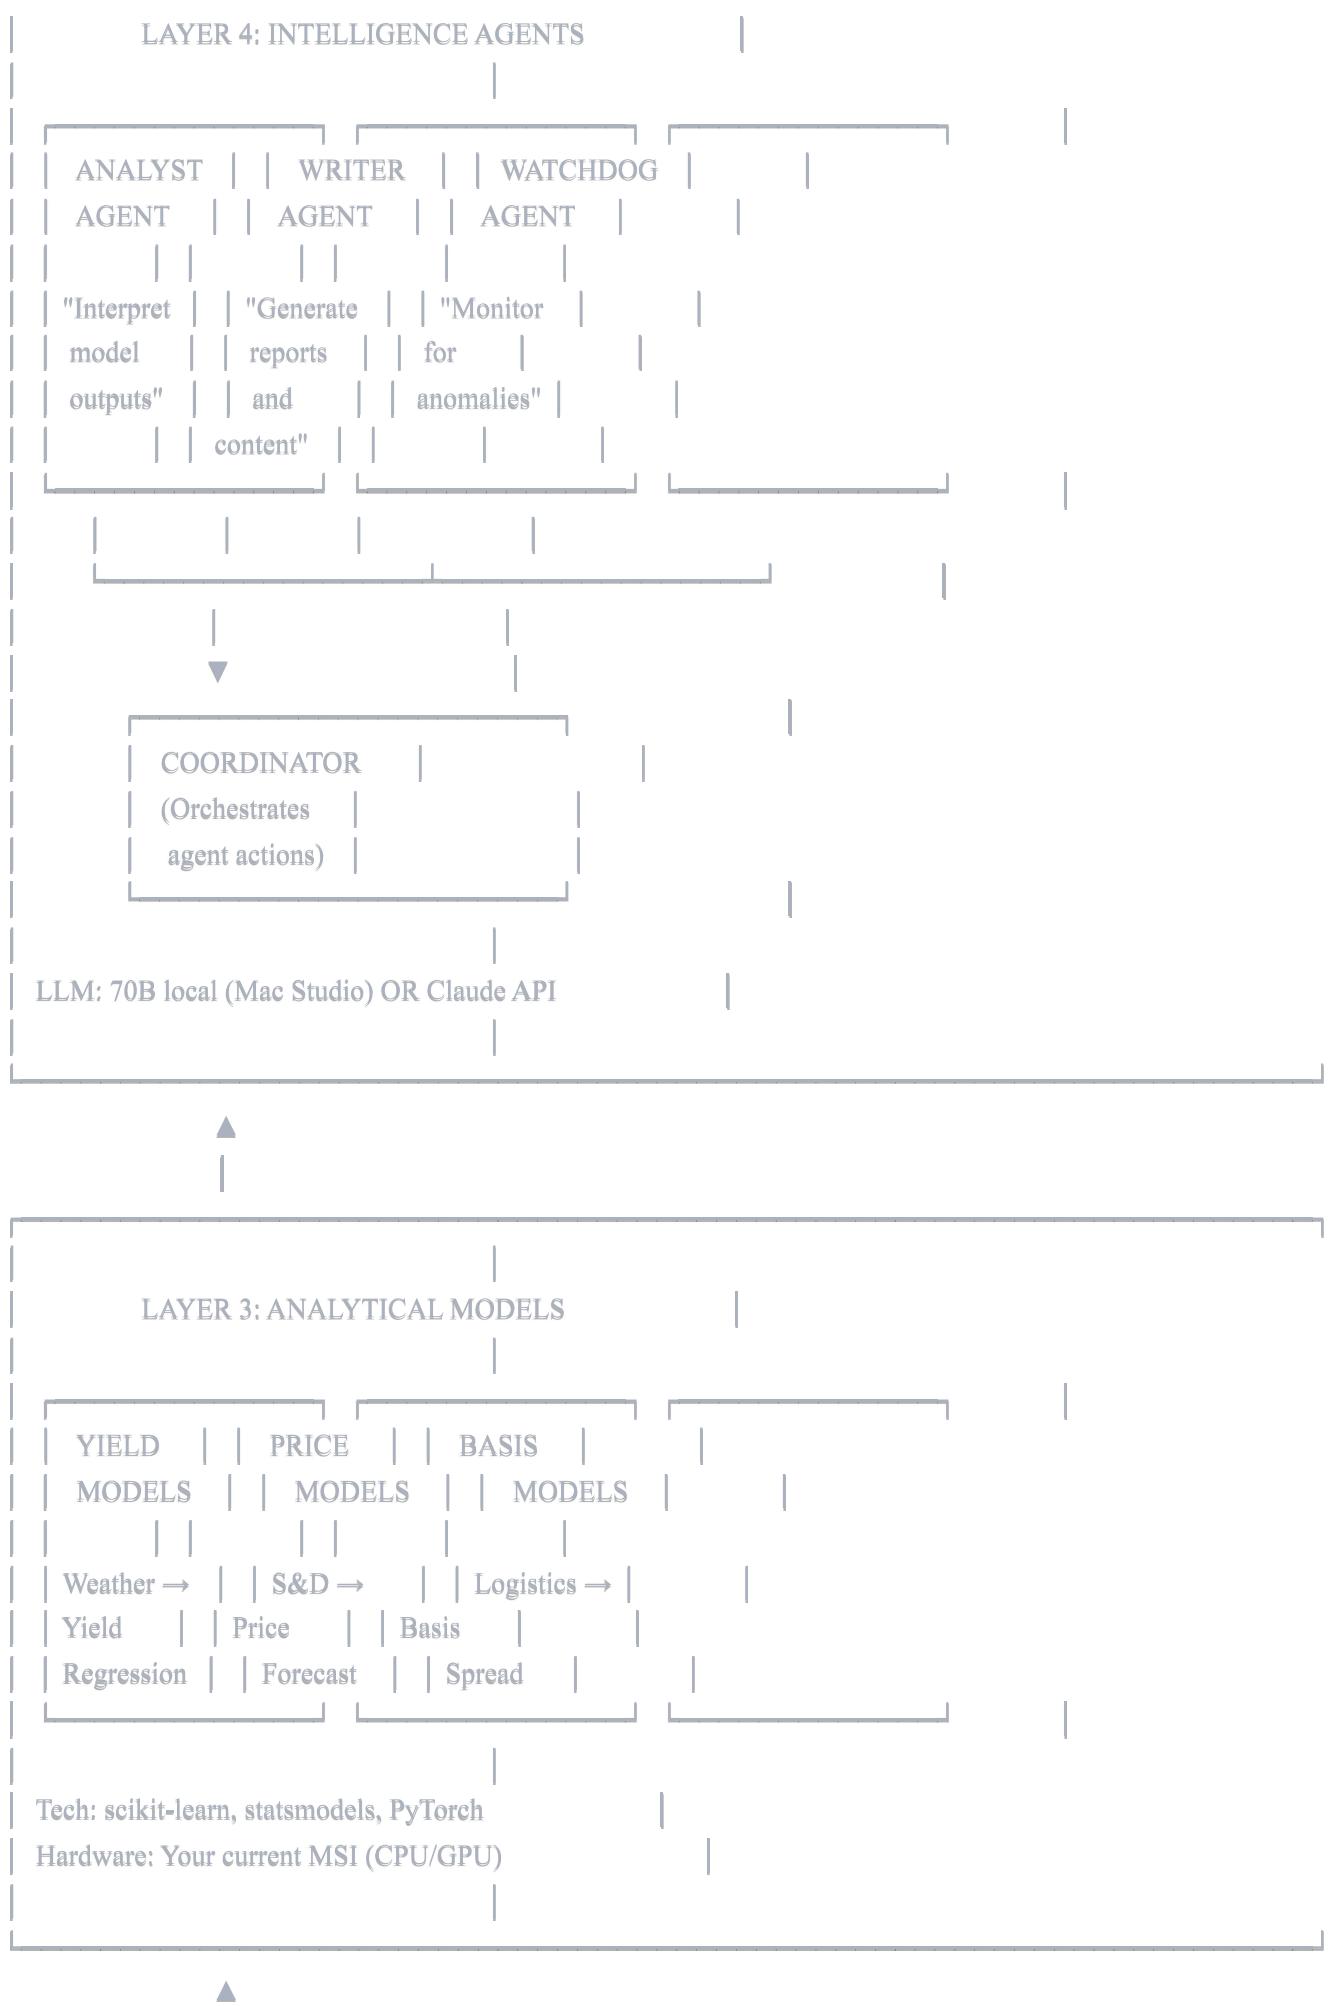
### Limitations:

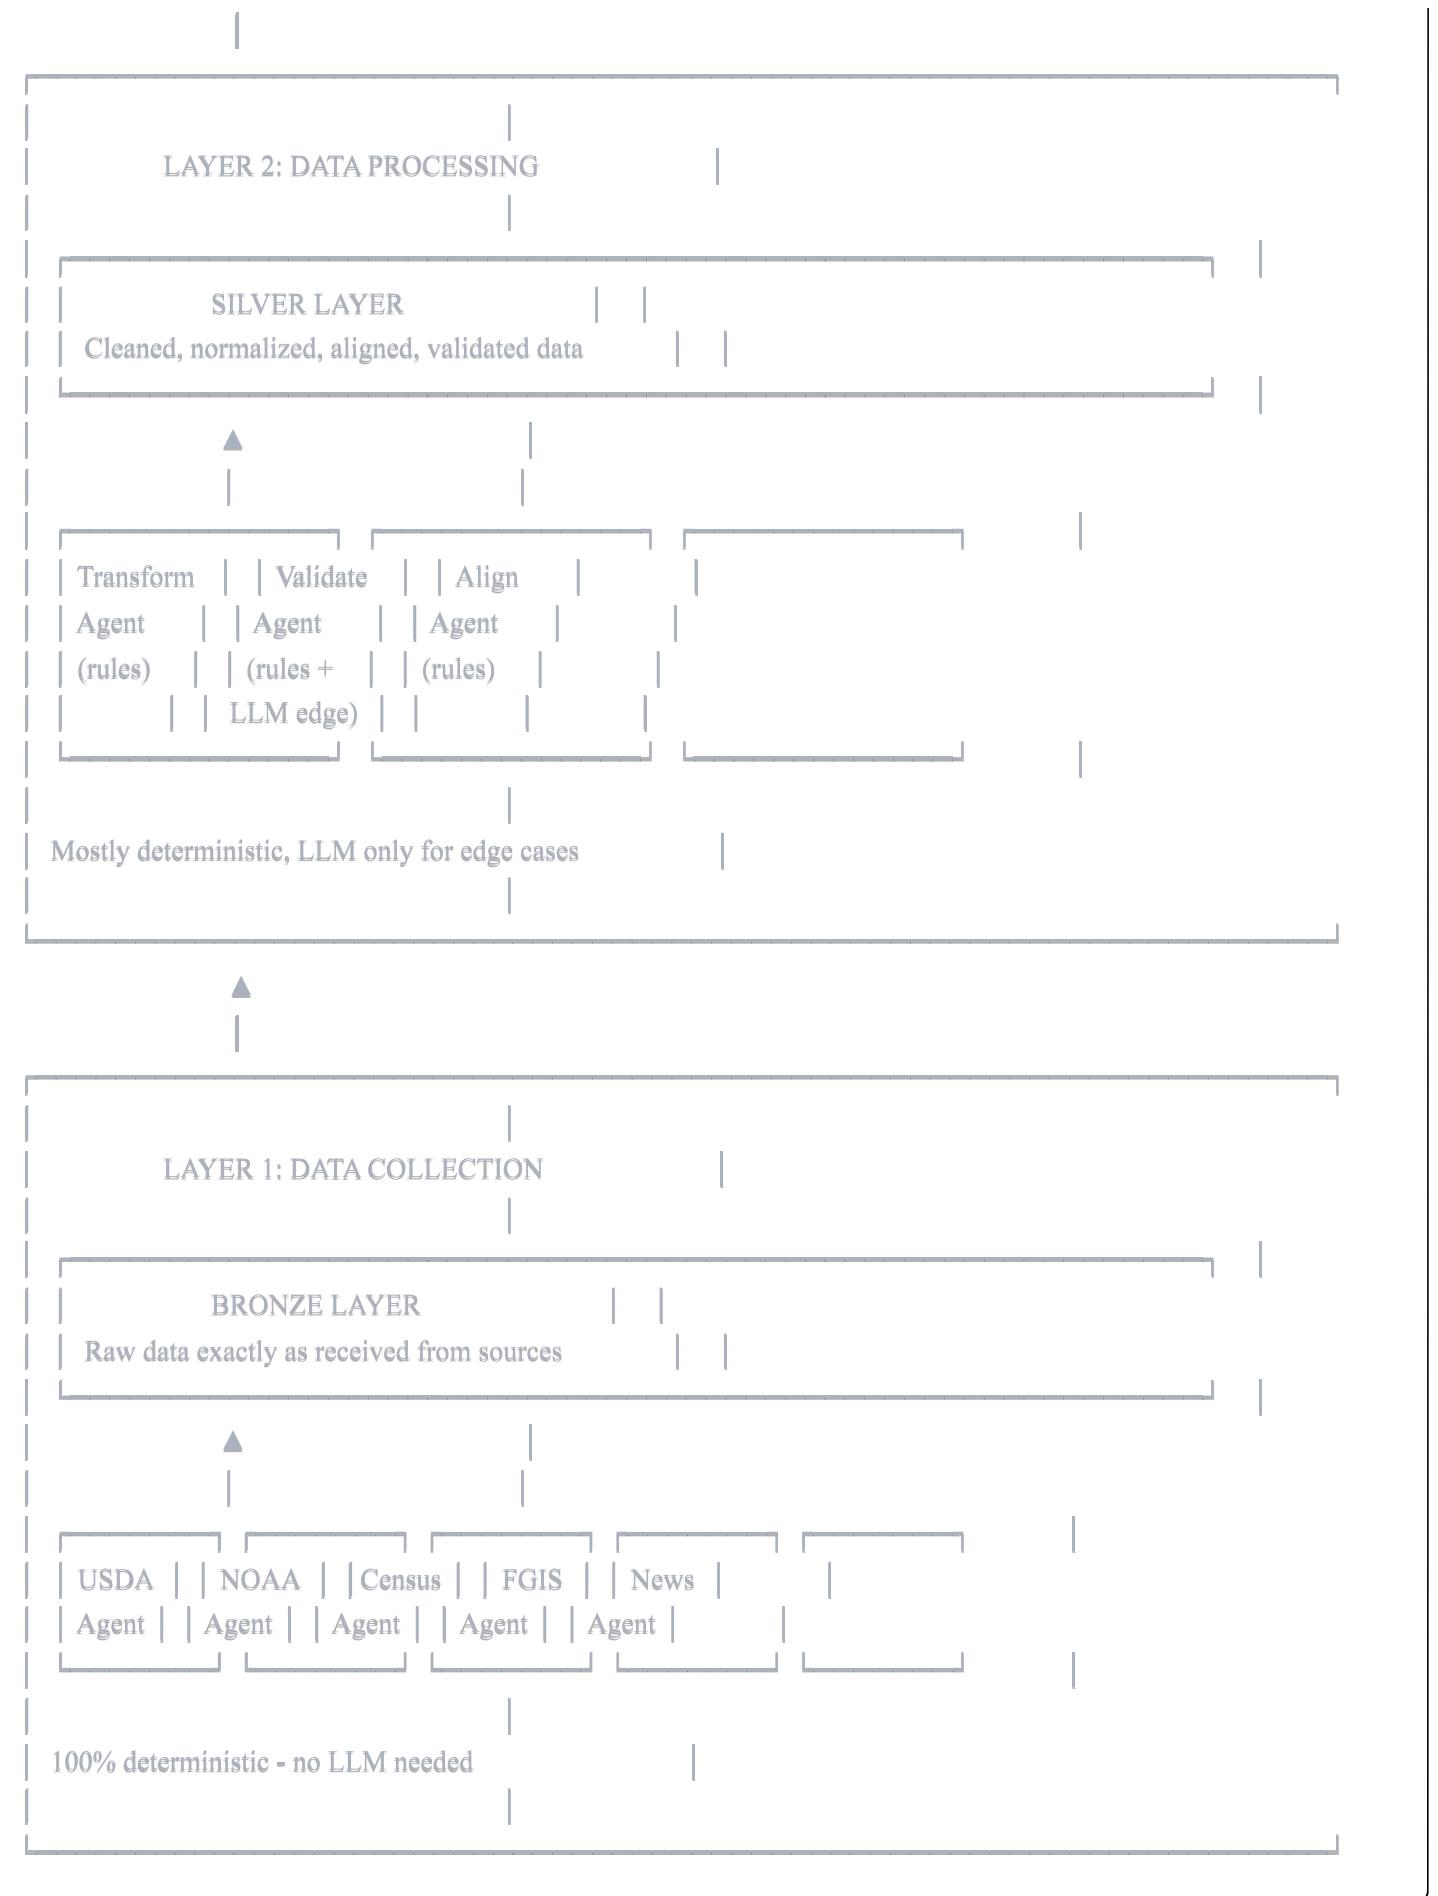
- Cannot handle unexpected situations
- Cannot synthesize across data sources
- Cannot explain "why" something is happening
- Cannot adapt without code changes
- Cannot generate nuanced analysis

## Part 2: Target Architecture

### The Layered Hybrid Model

LAYER 5: HUMAN INTERFACE  
Dashboards, Reports, Alerts, Chat Interface





## Intelligence Agent Specifications

### ANALYST AGENT

yaml

**name:** Analyst Agent

**goal:** Interpret quantitative model outputs and market data to form market views

**tools:**

- **query\_database:** Access Gold layer data
- **get\_model\_output:** Retrieve latest forecasts
- **get\_historical\_context:** Pull comparison periods
- **get\_news\_summary:** Recent relevant news
- **update\_market\_view:** Store current thesis

**triggers:**

- New model output available
- Significant data revision
- Price moves beyond threshold
- Scheduled daily review

**example\_reasoning:** |

"The yield model updated with new NOAA data showing 2" below-normal precip for western Iowa through mid-August. Model now projects state yield at 195 bu/acre vs 198 last week. This is a 1.5% reduction. Combined with similar revisions in Illinois, national production estimate drops ~150M bushels. Current futures appear to have priced in ~100M of this. Recommend: monitor for further weather deterioration; if precip deficit persists, market has room to rally \$0.15-0.20."

**llm\_requirement:** 70B+ for nuanced reasoning, or Claude API

### WRITER AGENT

yaml

**name:** Writer Agent

**goal:** Generate market reports and content in RLC's voice and style

**tools:**

- **get\_market\_view:** Current analyst thesis
- **get\_data\_tables:** Formatted S&D tables
- **get\_template:** Report templates
- **get\_style\_guide:** RLC writing standards
- **publish\_report:** Output to destination

**triggers:**

- Scheduled report times (daily, weekly)
- Analyst agent flags significant development
- Manual request

**output\_types:**

- Daily market brief (300-500 words)
- Weekly S&D summary (1000-1500 words)
- Flash alerts (50-100 words)
- Social media posts (LinkedIn, Twitter)
- Client emails

**llm\_requirement:** 7B adequate for templated content, 70B for original analysis

## WATCHDOG AGENT

yaml

**name:** Watchdog Agent

**goal:** Continuously monitor for anomalies, errors, and significant changes

**tools:**

- **query\_database:** Check latest data
- **compare\_historical:** Statistical anomaly detection
- **check\_data\_quality:** Validation rules
- **send\_alert:** Notify on findings
- **log\_observation:** Record for review

**monitors:**

- Data quality issues (missing, invalid, outliers)
- Unexpected price moves ( $> 2$  std dev)
- Unusual basis behavior
- Model forecast changes ( $>$  threshold)
- News keyword triggers
- Source availability (API health)

**alert\_levels:**

- **INFO:** Log only, include in daily digest
- **WATCH:** Include in next analyst review
- **ALERT:** Immediate notification
- **CRITICAL:** Interrupt current work

**ilm\_requirement:** 7B adequate for pattern matching, 70B for interpretation

## Part 3: Migration Phases

### Phase 0: Foundation (Current → Week 4)

**Goal:** Organize existing infrastructure, no new AI yet

**Tasks:**

1. Document all 46 current agents (inputs, outputs, dependencies)
2. Ensure Bronze/Silver/Gold data layers are clearly separated
3. Create comprehensive logging for all agent activities
4. Set up monitoring dashboard for system health
5. Identify which agents are truly needed vs. legacy

**Hardware:** Current MSI only

**LLM:** None (or existing small model usage)

**Cost:** \$0 (time investment only)

#### **Success Criteria:**

- Complete agent inventory document
  - All agents logging to central location
  - System health dashboard operational
  - Clear data lineage documentation
- 

### **Phase 1: Analytical Models (Weeks 4-12)**

**Goal:** Build the quantitative foundation that AI will interpret

#### **Tasks:**

1. Implement yield forecasting models
  - Weather → Yield regression (corn, soybeans by state)
  - Input: NOAA data, historical yields
  - Output: Point forecast + confidence interval
2. Implement price forecasting models
  - S&D → Price relationships
  - Seasonal patterns
  - Basis models by location
3. Create model output tables in Gold layer
  - Standardized format for all model outputs
  - Version tracking (what changed, when)
  - Forecast vs. actual tracking
4. Build model performance monitoring
  - Track forecast accuracy over time
  - Identify when models need retraining

**Hardware:** Current MSI (CPU for training, GPU optional)

**LLM:** None yet

**Cost:** \$0 (time investment only)

## **Success Criteria:**

- At least 3 working forecast models
- Model outputs flowing to Gold layer
- Accuracy tracking operational
- Documentation of model methodology

## **Key Deliverable:**

When the ANALYST AGENT later asks "what does the model say about Illinois soybean yields?", there's a real model with a real answer—not an LLM hallucination.

---

## **Phase 2: LLM Integration - Read Only (Weeks 12-16)**

**Goal:** Add LLM reasoning that observes but doesn't act

### **Tasks:**

1. Set up LLM inference
  - Option A: Claude API (recommended for this phase)
  - Option B: Local 7B/13B on RTX 5080 via Ollama
2. Create WATCHDOG AGENT (read-only)
  - Monitors data quality
  - Flags anomalies
  - Generates daily digest
  - NO automated actions, just observations
3. Create prototype ANALYST AGENT (advisory only)
  - Given: Latest model outputs + recent data
  - Produces: Draft market interpretation
  - Human reviews before any use
4. Evaluation period
  - Track: How often is the LLM's judgment correct?
  - Track: How often does it catch things you missed?
  - Track: How often does it hallucinate or miss obvious issues?

**Hardware:** Current MSI + API access

**LLM:** Claude API (~\$50-150/month estimated)

**Cost:** ~\$200-600 total for phase

#### Success Criteria:

- WATCHDOG running daily, generating useful observations
- ANALYST producing draft interpretations you find valuable
- Quantified accuracy: "LLM was helpful X% of the time"
- Clear understanding of LLM limitations in your domain

#### Decision Point:

At the end of Phase 2, you'll know:

- Is the LLM reasoning valuable enough to continue?
- Is 7B/13B local sufficient, or do you need 70B?
- Are API costs acceptable, or should you go local?

If LLM adds little value → Stay with Phase 1 architecture

If LLM is valuable + API costs OK → Continue with API

If LLM is valuable + want local → Consider hardware purchase

### Phase 3: LLM Integration - Write Access (Weeks 16-24)

**Goal:** Allow AI agents to take limited automated actions

#### Tasks:

1. WRITER AGENT implementation
  - Generates draft reports from templates + model outputs
  - Human approval required before publishing
  - Gradually reduce approval requirements for routine content
2. WATCHDOG AGENT upgrades
  - Can auto-log issues to tracking system
  - Can trigger data re-fetch on quality failures
  - Still requires human approval for alerts to you
3. Implement agent coordination
  - WATCHDOG detects issue → triggers ANALYST review

- ANALYST forms view → triggers WRITER for report
  - Simple workflow, human checkpoints
4. Build feedback loop
    - When you correct agent output, log the correction
    - Use corrections to improve prompts/fine-tuning

**Hardware:** Current MSI + API (or local if decided in Phase 2)

**LLM:** Same as Phase 2 decision

**Cost:** Depends on Phase 2 decision

#### **Success Criteria:**

- WRITER producing usable first drafts 70%+ of the time
  - WATCHDOG catching issues before you notice them
  - Clear escalation paths working
  - Reduced time spent on routine analysis
- 

### **Phase 4: Autonomous Operations (Weeks 24+)**

**Goal:** Agents operate with minimal supervision for routine tasks

#### **Tasks:**

1. Define "routine" vs "exceptional" clearly
  - Routine: Daily data refresh, standard reports, normal market conditions
  - Exceptional: Data anomalies, unusual market moves, breaking news
2. Implement tiered autonomy
  - TIER 1 (Full auto): Data collection, validation, standard reports
  - TIER 2 (Log + proceed): Minor anomalies, routine analysis updates
  - TIER 3 (Ask first): Significant findings, non-routine content
  - TIER 4 (Human only): Trading signals, client communications, model changes
3. Implement COORDINATOR agent
  - Orchestrates other agents
  - Decides which agent to invoke based on situation
  - Manages priorities and resources
4. Continuous improvement

- Agents suggest improvements to their own prompts
- You approve/reject, system learns

**Hardware:** Possibly Mac Studio if local 70B needed

**LLM:** 70B local or continued API

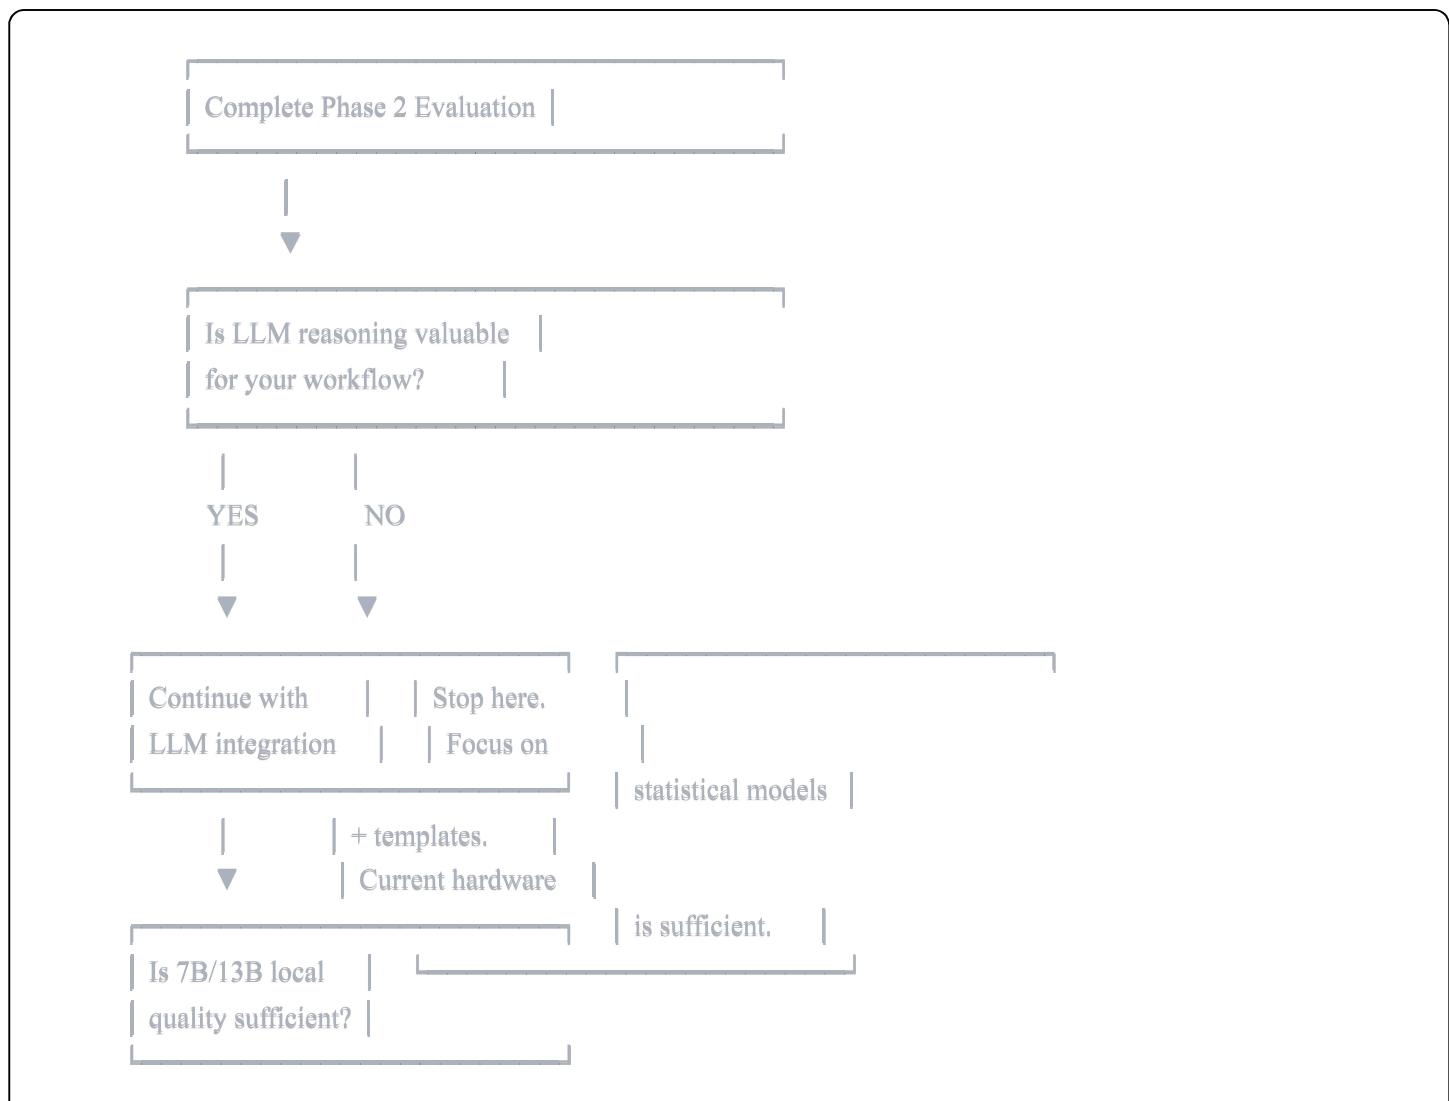
**Cost:** \$0-6,000 depending on hardware decision

#### Success Criteria:

- System runs for days without intervention (for routine operations)
  - You're notified only of things that need your attention
  - Report quality indistinguishable from your manual work
  - Clear audit trail of all agent decisions
- 

### Part 4: Hardware Decision Framework

#### Decision Tree



YES

NO

Stay with current MSI + Ollama  
Need larger model (70B+)  
Cost: \$0

Are API costs acceptable long-term?  
(~\$100-300/month)

YES

NO

Continue with Claude/GPT API  
Purchase Mac Studio M2/M3 Ultra 192GB  
Cost: ~\$1,200-3,600/yr | Cost: \$2,000-6,500 (one-time)

## Hardware Options Summary

Option	Upfront Cost	Monthly Cost	Model Size	Best For
Current MSI + 7B Ollama	\$0	\$0	7B-13B	Simple tasks, testing
Current MSI + Claude API	\$0	\$100-300	Frontier	Best quality, low volume
Mac Studio M2 Ultra 192GB (used)	\$2,000-2,500	\$0	70B	Good quality, high volume, budget
Mac Studio M3 Ultra 192GB (new)	\$6,000-6,500	\$0	70B	Good quality, high volume, warranty
Mac Studio M3 Ultra 512GB (new)	\$10,000-14,000	\$0	180B+	Maximum future-proofing

---

## Part 5: Specific Agent Designs

### Example: Southern Illinois Soybean Yield Agent

This addresses your specific example: "forecast the yield of soybeans in southern Illinois based on weather/yield models and write an analysis"

```
yaml
```

**name:** IL\_Soybean\_Yield\_Agent

**type:** Composite (Statistical Model + LLM Interpretation)

**components:**

**1\_data\_collection:**

**type:** Deterministic Script

**schedule:** Daily at 6am CT

**actions:**

- Fetch NOAA weather data for IL climate divisions 6, 7, 8, 9
- Fetch soil moisture from NLDAS
- Fetch crop condition from NASS (when available)
- Store in Bronze layer

**2\_model\_execution:**

**type:** Statistical Model (Python)

**schedule:** After data collection completes

**model\_type:** Ridge regression with weather features

**features:**

- Cumulative precipitation (Apr-Aug)
- GDD accumulation
- Days above 95°F in July-Aug
- Soil moisture at key growth stages
- Trend yield adjustment

**output:**

- Point estimate (bu/acre)
- 80% confidence interval
- Comparison to 5-year average
- Comparison to last week's estimate

**store:** Gold layer (model\_outputs.il\_soybean\_yield)

**3\_interpretation:**

**type:** LLM Agent

**trigger:** Model output changes by > 1 bu/acre OR weekly schedule

**prompt:** |

You are an agricultural analyst specializing in Midwest soybean production.

**Current model output for Southern Illinois soybeans:**

- **Yield estimate:** {yield\_estimate} bu/acre
- **Last week:** {last\_week\_estimate} bu/acre
- **Change:** {change} bu/acre ({change\_pct}%)
- **5-year average:** {avg\_5yr} bu/acre
- **Confidence interval:** {ci\_low} to {ci\_high}

**Weather factors driving change:**

{weather\_summary}

**Recent observations:**

{crop\_condition\_notes}

**Tasks:**

1. Explain what's driving the yield estimate change in plain language
2. Assess whether the model's reaction seems appropriate given conditions
3. Note any factors the model might be missing
4. Provide a 2-3 sentence summary suitable for a client report

**output:** Structured interpretation stored to Gold layer

**4\_content\_generation:**

**type:** LLM Agent (WRITER)

**trigger:** Interpretation complete + scheduled report time

**inputs:**

- Interpretation from step 3
- **Template:** weekly\_crop\_update.md
- **Style guide:** rlc\_writing\_standards.md
- **Historical examples:** sample\_crop\_reports/

**output:**

- Draft report section (200-400 words)
- Suggested charts/tables
- Confidence flag (high/medium/low)

**review:** Human approval required before publishing

## Example: Data Quality Watchdog

yaml

**name:** Data\_Quality\_Watchdog

**type:** Hybrid (Rules + LLM for ambiguous cases)

monitors:

**price\_data:**

**rules:**

- Price must be positive
- Daily change < 15% (configurable)
- No gaps ≥ 2 trading days
- Timestamp must be valid

**llm\_escalation:**

**trigger:** Rule violation detected

**prompt:** |

A price data anomaly was detected:

**Commodity:** {commodity}

**Date:** {date}

**Reported price:** {price}

**Previous price:** {prev\_price}

**Change:** {change\_pct}%

**Historical context:**

- **30-day average:** {avg\_30d}
- **30-day std dev:** {std\_30d}
- **Recent news:** {news\_summary}

**Is this likely:**

- A) A data error that should be flagged for manual review
- B) A real market move that should be accepted
- C) Unclear - need more information

Explain your reasoning.

**fundamental\_data:**

**rules:**

- Production values within historical range ( $\pm 50\%$ )
- Stocks cannot be negative
- Exports + domestic use  $\leq$  total supply
- Marketing year dates must align

**llm\_escalation:**

**trigger:** Rule violation OR data revision > 5%

**prompt:** |

A fundamental data change was detected:

Report: {report\_name}

Field: {field\_name}

Previous value: {old\_value}

New value: {new\_value}

Change: {change\_pct}%

This change is {larger/smaller} than typical revisions.

Historical revision patterns for this field:

{revision\_history}

Is this revision:

- A) Within normal bounds - accept automatically
- B) Unusual but plausible - log for review
- C) Likely an error - flag for immediate attention

Explain your reasoning.

alert\_routing:

A\_accept: Log only, continue processing

B\_review: Add to daily digest, continue processing

C\_flag: Send immediate alert, pause downstream processing

## Part 6: Risk Assessment

### Things That Could Go Wrong

Risk	Likelihood	Impact	Mitigation
LLM hallucinates market data	Medium	High	Always ground in database queries, never trust LLM for numbers
LLM gives confidently wrong analysis	Medium	Medium	Human review for all external-facing content
Model drift (forecasts degrade)	High	Medium	Continuous accuracy tracking, scheduled retraining
Over-reliance on AI judgment	Medium	High	Maintain your own market view, use AI as input not oracle

Risk	Likelihood	Impact	Mitigation
API costs exceed budget	Low	Low	Set spending caps, monitor usage
Hardware failure during critical period	Low	High	Redundancy plan, cloud backup capability
Data source changes format/breaks	High	Medium	Robust error handling, multiple source validation

## The Honest Limitations

### What this architecture CAN do:

- Automate routine data processing and validation
- Generate first drafts of reports and analysis
- Catch anomalies and flag them for your attention
- Apply your analytical frameworks consistently
- Work 24/7 without fatigue

### What this architecture CANNOT do:

- Predict prices better than your statistical models
- Replace your market intuition and experience
- Handle truly novel situations without guidance
- Make trading decisions you'd trust without review
- Understand context it hasn't been taught

### The 1+1=3 you're looking for:

- Your expertise + AI synthesis = faster, more comprehensive analysis
- Statistical models + LLM interpretation = numbers with narrative
- Continuous monitoring + intelligent filtering = you see only what matters
- Templated output + LLM flexibility = consistent quality, less drudgery

---

## Part 7: Success Metrics

### Phase 1 Success (Statistical Models)

- 3+ forecast models operational

- Model accuracy tracked and documented
- Gold layer contains all model outputs

### **Phase 2 Success (LLM Read-Only)**

- WATCHDOG catching issues before manual discovery (50%+ of issues)
- ANALYST interpretations rated "useful" 70%+ of time
- Clear quantification of LLM value-add

### **Phase 3 Success (LLM Write Access)**

- WRITER drafts accepted with minor edits 60%+ of time
- Time spent on routine reports reduced by 50%+
- No significant errors in published content

### **Phase 4 Success (Autonomous Operations)**

- System runs 7+ days without intervention (routine operations)
  - Alert signal-to-noise ratio > 3:1 (useful alerts vs. false alarms)
  - Your weekly time on routine analysis reduced by 70%+
- 

## **Part 8: Next Steps**

### **Immediate Actions (This Week)**

- Review and refine this architecture document
- Complete inventory of current 46 agents
- Identify which current agents map to which layer
- Set up centralized logging if not already done

### **Phase 0 Deliverables (Next 4 Weeks)**

- Complete agent documentation
- System health dashboard
- Data lineage documentation
- Decision: Which statistical models to build first?

### **Key Decision Points**

- **Week 4:** Ready to start Phase 1? (statistical models)
- **Week 12:** Ready to start Phase 2? (LLM integration)
- **Week 16:** LLM value assessment → hardware decision

- **Week 24:** Autonomous operations readiness assessment
- 

## Appendix A: Technology Stack

### Recommended Stack

Layer	Technology	Notes
Data Collection	Python + requests/aiohttp	Keep existing
Data Storage	PostgreSQL	Keep existing
Data Processing	Python + pandas	Keep existing
Statistical Models	scikit-learn, statsmodels	Add
LLM Inference (local)	Ollama + llama.cpp	Add
LLM Inference (API)	Claude API / OpenAI API	Add
Agent Framework	LangChain or custom	Evaluate
Orchestration	Prefect or existing scheduler	Evaluate
Monitoring	Grafana or custom dashboard	Add
Version Control	Git	Keep existing

### File Structure (Proposed)

```

rlc-agent/
  ├── agents/
  │   ├── collection/    # Layer 1: Data fetching
  │   ├── processing/   # Layer 2: Transformation
  │   ├── models/       # Layer 3: Statistical models
  │   └── intelligence/ # Layer 4: LLM agents
  └── data/
      ├── bronze/      # Raw data
      ├── silver/      # Cleaned data
      └── gold/        # Analytical outputs
  └── config/
      ├── agents.yaml  # Agent configurations
      └── models.yaml  # Model parameters

```

```
|   └── prompts/    # LLM prompt templates  
|   ├── templates/  
|   |   └── reports/    # Report templates  
|   |   ├── monitoring/  
|   |   |   └── dashboards/    # Monitoring configs  
|   |   └── docs/  
|   |       └── architecture.md # This document
```

---

## Appendix B: Glossary

- **Agent:** A software component that performs tasks, ranging from simple scripts to AI-powered decision makers
  - **Bronze Layer:** Raw data stored exactly as received from sources
  - **Silver Layer:** Cleaned, validated, normalized data
  - **Gold Layer:** Analytical outputs, model results, aggregations
  - **LLM:** Large Language Model (e.g., Claude, GPT-4, Llama)
  - **Frontier Model:** State-of-the-art LLM (Claude, GPT-4) with best reasoning capability
  - **Local Model:** LLM running on your own hardware
  - **Agentic:** Software that can decide what actions to take based on goals, not just explicit instructions
- 

*Document Status: Draft for Review Last Updated: February 4, 2026*